


AI Memory Problem: Why Bigger Context Windows Don't Stop Agents From Forgetting

Kabui, Charles

2026-06-03

[Read at ToKnow.ai](#)



Even with million-token context windows in Claude and Gemini, AI agents forget everything between sessions and miss information buried in long inputs. A [Stanford study](#) confirmed that retrieval degrades sharply when relevant information sits in the middle of long contexts, and bigger windows alone don't fix this. Three open-source projects attack the problem differently. [EverMemOS](#) converts conversations into structured memory cells, consolidates them

into thematic scenes, and achieves state-of-the-art on the [LoCoMo](#) and [LongMemEval](#) benchmarks. [Supermemory](#) ranks #1 on all three major memory benchmarks (LongMemEval at 81.6%, LoCoMo, and ConvoMem) with automatic fact extraction and contradiction resolution. [Agentmemory](#) targets coding agents: 12 auto-capture hooks record tool usage, compress sessions into a 4-tier memory hierarchy (working, episodic, semantic, procedural), and report 92% fewer tokens with 95.2% retrieval recall.

The cost adds up. At Claude Opus 4.8's \$5/\$25 per million token pricing, a developer running 50 sessions a day can spend \$50-100 just re-reading files the agent already processed. Agentmemory cuts session context to roughly 1,900 tokens from 22,000+. Supermemory's user profiles return in about 50ms. [LongTraceRL](#) from Tsinghua attacks from the model side, using reinforcement learning with rubric rewards to train 4B-30B parameter models to navigate long contexts more reliably across five benchmarks.

The context problem is splitting into two engineering disciplines. Longer windows and [KV cache compression](#) solve input capacity. Persistent memory systems solve continuity. The memory layer is becoming its own infrastructure stack.

Sources:

- [Lost in the Middle: How Language Models Use Long Contexts \(Stanford\)](#)
- [EverMemOS: Self-Organizing Memory Operating System \(arXiv\)](#)
- [Supermemory: Memory Engine for AI](#)
- [Agentmemory: Persistent Memory for AI Coding Agents](#)
- [LongTraceRL: Long-Context Reasoning from Search Trajectories \(Tsinghua\)](#)

***Disclaimer:** For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)*