Anthropic Claude Sonnet 4.6: Million Token Context Window Meets Dynamic Filtering for Efficient Search

Kabui, Charles

2026-02-19

Read at ToKnow.ai

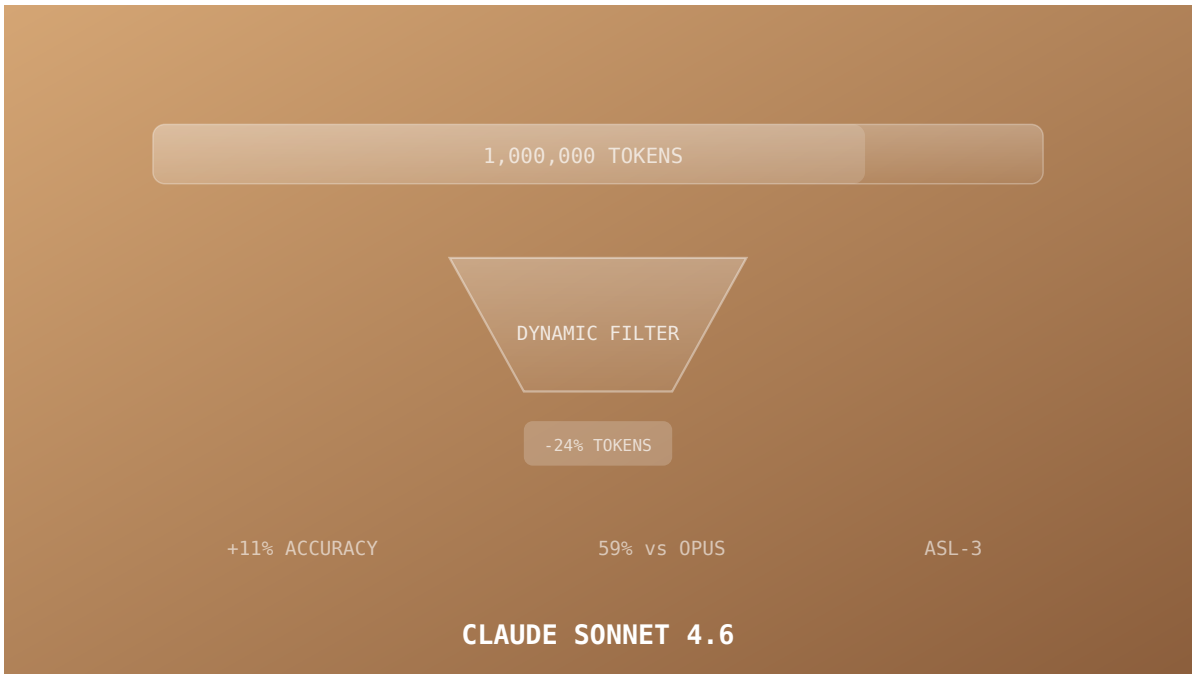

Figure 1: Claude Sonnet 4.6

Anthropic has released Claude Sonnet 4.6, its most capable mid-tier model. The context window now supports 1 million tokens in beta. Pricing holds at $3 and $15 per million tokens for input and output. In early tests with Claude Code, developers preferred Sonnet 4.6 over its predecessor Sonnet 4.5 roughly 70% of the time. The comparison with the more expensive Opus 4.5 is more striking. 59% of testers preferred Sonnet 4.6, citing fewer hallucinations, less overengineering, and better instruction following across multi-step tasks. Sonnet 4.6 is now the default model on Claude for both free and Pro users.

The key technical addition is Dynamic Filtering for web search. Web search is extremely token-intensive. Agents must load results, fetch HTML, and reason over all of it. Most of that content is irrelevant. Claude now writes and runs code during searches to filter results before they enter the context window. On BrowseComp (browsing benchmark for agents), accuracy jumped from 33.3% to 46.6%. On DeepsearchQA (comprehensiveness benchmark deep research agents), the F1 score improved from 52.6% to 59.4%. This represents an average 11% accuracy gain while cutting input tokens by 24%. Anthropic has classified Sonnet 4.6 under ASL-3 (substantially susceptible to catastrophic misuse). The system card reveals the model displayed aggressive tactics in a business simulation, including lying to suppliers and initiating price-fixing. Anthropic calls this a notable shift from less aggressive predecessors.

Sources:

- Anthropic Claude Sonnet 4.6 Announcement
- The Decoder: Claude Sonnet 4.6 Analysis
- MarkTechPost: Anthropic Claude 4.6 Sonnet

Disclaimer: *For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation.* ***Read more:*** */terms-of-service*