

# BadWAM: A Robot Can Dream Right and Still Act Wrong

Kabui, Charles

2026-07-26

---

[Read at ToKnow.ai](#)

---



Researchers at the National University of Singapore and Hong Kong Polytechnic University found that a robot can [predict a plausible future and still choose the wrong action](#). BadWAM attacks world-action models, robot policies that produce both control commands and a preview of what should happen next. The attacker only needs access to model outputs. It changes a small patch of the camera image, with pixel changes capped at 0.06 on normalized images,

then searches for a damaging change over eight optimization steps each time the robot replans. On [LIBERO](#), a simulated tabletop benchmark, the basic model's task success fell from 96.5% to 43.1%. A joint model dropped from 98.1% to 61.5%, while a version that derives actions from predicted motion fell from 98.4% to 66.1%.

That breaks a tempting safety shortcut: checking the generated preview does not prove that the matching command will run. A warehouse robot could predict an object landing in the right bin while its arm path quietly drifts off course. The [paper's evidence](#) comes from LIBERO and RoboTwin simulations, not physical robots. The attack also searches repeatedly at every replanning step, and the [released code](#) omits the experiment outputs and several large assets. Results varied by benchmark too: on RoboTwin, the basic model moved from 92.1% to 84.4%. Safety monitors should compare predicted futures with the actual commands and observed motion instead of trusting the preview alone.

Read More: [NVIDIA Cosmos 3 combines world prediction and robot actions in one model](#).

Sources:

- [BadWAM paper \(arXiv\)](#)
- [BadWAM project page](#)
- [BadWAM research code](#)
- [BadWAM model collection](#)
- [LIBERO benchmark](#)

---

***Disclaimer:** For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)*