

Smart Web Crawling: A Dynamic and Controllable Approach to Web Crawling

Crawling Kenyan Legal Documents with Python

Kabui, Charles

2025-08-03

Table of Contents

The Challenge: Precision Over Scale	2
State Management with SQLite	3
Dynamic JavaScript Execution	4
Real-World Application: KenyaLaw.org	4
The Problem	4
The Solution	4
Google Colab Integration	5
Controlled Resource Management	5
When to Use This Approach	5
Small Scale Crawling Advantage	6
Code	6
Colab Setup	6
Import Required Libraries	7
Browser Functions	7
Sqlite Database Manager	10
Schema	14
Initialize the Crawler	21
Conclusion	23

 [Read at ToKnow.ai](#)

 [Open in Kaggle](#)

 [Download as Notebook](#)




Figure 1: Smart Web Crawling: A Dynamic and Controllable Approach to Web Crawling

The Challenge: Precision Over Scale

While enterprise crawling solutions like [Scrapy](#) or [Apache Nutch](#) excel at massive-scale data extraction, small scale crawling need something different: **controlled, schema-driven crawling with surgical precision.**

This approach isn't about crawling millions of pages per hour. It's about crawling the *right* pages, in the *right* way, with complete control over the process.

Schema-Based Crawling allows you to define how your crawler interacts with different websites, making it adaptable to various site structures and behaviors.

Each target site is treated as a unique ecosystem with its own rules.

Here is an example of a schema for crawling <kenyalaw.org>:

```
const schemas = [
  {
    prepare_page_expression: () => {
      return new Promise((resolve) => setTimeout(resolve, 10 * 1000));
    },
    locations: [
      { regex: ".*kenyalaw.org.*", examples: ["https://kenyalaw.org/"] },
    ],
    ignorePages: [
      {
        regex: "^(?!.*kenyalaw\\.org).*",
        examples: ["https://www.facebook.com/MyKenyaLaw"],
      },
    ],
  },
];
```

This schema tells our crawler: - **Where to operate:** `locations` field specifies the domains/locations the crawler is permitted to explore, only on <kenyalaw.org> domains - **What to ignore:** `ignorePages` prevents the crawler from visiting undesirable links, such as social media pages - **How to prepare:** `prepare_page_expression` field allows for custom JavaScript to execute before link extraction

State Management with SQLite

A persistent database is crucial. This approach uses an SQLite database to store and manage the discovered links and their status.

```
class Status(Enum):
    DOWNLOADED = "downloaded"
    CRAWL = "crawl"
    CRAWLED = "crawled"
```

Each URL has a status, retry count, and timestamps. This enables: - **Resumable crawling:** Stop and restart without losing progress - **Retry logic:** Handle temporary failures gracefully - **Analytics:** Track crawling progress and patterns

Dynamic JavaScript Execution

The system generates JavaScript that: - Validates each domain against schema rules - Executes custom preparation code - Extracts and filters links based on ignore patterns - Returns clean, deduplicated URLs

```
const isLocationAllowed = schema.locations.some((domain) =>
  new RegExp(domain.regex).test(location.href)
);

if (isLocationAllowed) {
  if (schema.prepare_page_expression) {
    await Promise.resolve(schema.prepare_page_expression());
  }
  // Extract links...
}
```

Real-World Application: KenyaLaw.org

Let's see this in action with a legal database crawler:

The Problem

KenyaLaw.org contains thousands of legal documents across multiple formats: - Case law with metadata - PDF exports - XML structured data - Search result pages

The Solution

Our schema-based approach:

1. **Discovers** all case law URLs through search pages
2. **Ignores** direct external links, eg: social media links
3. **Handles** JavaScript-heavy search interfaces
4. **Tracks** progress in SQLite for reliability

```
# Start crawling from advanced search
await crawl_and_save("https://kenyalaw.org/caselaw/cases/advanced_search/")

# System automatically discovers:
# - Individual case pages
# - Additional search result pages
# - Related legal document links
# While ignoring export endpoints
```

Google Colab Integration

One unique advantage: this runs beautifully in Google Colab. Mount your Google Drive, point the database there, and you have persistent, resumable crawling with zero local setup:

```
from google.colab import drive
drive.mount("/content/drive")
database_path_str = "/content/drive/MyDrive/legal_crawl/database.db"
```

Your crawling state persists across Colab sessions, and you can analyze results directly in notebooks.

Controlled Resource Management

The system includes smart resource management:

```
refresh_browser_tabs = 50 # Restart browser every 50 tabs
crawl_count = 10          # Process 10 URLs per batch
```

This prevents memory leaks and browser crashes during long-running crawls while maintaining reasonable throughput.

When to Use This Approach

Perfect for: - Legal document databases - Academic research sites - News archives with complex navigation - Sites requiring authentication or session management - JavaScript-heavy applications - Projects needing precise control over crawling behavior

Not suitable for: - High-volume, simple website crawling - Basic e-commerce product scraping - Sites with robust APIs - When speed trumps control

Small Scale Crawling Advantage

This approach shines when you need:

- **Precision over speed:** Quality links, not quantity
- **Flexibility:** Easy schema updates for new sites
- **Observability:** Full visibility into crawling progress
- **Resumability:** Handle interruptions gracefully
- **Cost control:** Run on free Google Colab instead of expensive cloud crawling services

Code

```
# Path to the SQLite database file
database_path_str = "./_database.db"

# Number of browser tabs to open before refreshing the browser (prevents memory leaks)
refresh_browser_tabs = 50

# Number of URLs to process per crawling batch
crawl_count = 10
```

Colab Setup

If you are running this in Google Colab, uncomment and run the following code to set up the environment:

```
# # Install required Python packages and Playwright browser binaries
# !pip install playwright peewee jsbeautifier
# !playwright install

# # Mount Google Drive to access persistent storage in Colab
# from google.colab import drive
# drive.mount("/content/drive")

# # Set the database path to a location on Google Drive for persistence
# database_path_str = "/content/drive/MyDrive/legal_crawl/database.db"
```

Import Required Libraries

```
import asyncio
import json
import shutil
from datetime import datetime, timezone
from enum import Enum
from pathlib import Path
from typing import Any, Callable, Coroutine, List, Literal, Tuple, TypedDict

import jsbeautifier
import peewee
from playwright.async_api import (
    Browser,
    Page,
    Playwright,
    async_playwright,
)
```

Browser Functions

```
async def get_browser_async(
    p: Playwright, executable_path: str = None, *, headless=True, incognito=False
) -> Browser:
    """
    Launches a Chromium browser instance using Playwright.

    Args:
        p (Playwright): The Playwright instance.
        executable_path (str, optional): Path to the Chromium/Chrome executable.
            If None, tries to auto-detect.
        headless (bool, optional): Whether to run browser in headless mode. Default is True.
        incognito (bool, optional): Whether to launch browser in incognito mode. Default is False.

    Returns:
        Browser: The launched Playwright browser instance.
    """
    executable_path = (
        executable_path
        or shutil.which("google-chrome")
    )
```

```

        or shutil.which("chromium-browser")
        or shutil.which("chromium")
        or None
    )

    browser = await p.chromium.launch(
        headless=headless,
        executable_path=executable_path,
        args=["--disable-web-security"] + (["--incognito"] if incognito else []),
    )
    return browser

async def get_browser_page_async(
    executable_path: str = None,
    width: int = 0,
    height: int = 0,
    *,
    playwright: Playwright = None,
    headless: bool = True,
    incognito=False,
    refresh_browser_tabs: int = None
) -> tuple[Page, Callable[[], Coroutine[Any, Any, None]]]:
    """
    Returns a factory for new Playwright pages and a cleanup function.

    Args:
        executable_path (str, optional): Path to the Chromium/Chrome executable.
        width (int, optional): Viewport width for new pages. If 0, uses default.
        height (int, optional): Viewport height for new pages. If 0, uses default.
        playwright (Playwright, optional): Playwright instance. If None, starts a new one.
        headless (bool, optional): Whether to run browser in headless mode. Default is True.
        incognito (bool, optional): Whether to launch browser in incognito mode. Default is False.
        refresh_browser_tabs (int, optional): Number of tabs to open before refreshing the browser.

    Returns:
        tuple: (new_page function, close_playwright function)
            - new_page(): Coroutine that returns a new Playwright Page.
            - close_playwright(): Coroutine to close browser and Playwright resources.
    """
    playwright = await async_playwright().start()
    browser = None

```

```

context = None
tabs_created = 0

async def new_page() -> Page:
    nonlocal browser
    nonlocal context
    if (
        browser == None
        or context == None
        or (refresh_browser_tabs and tabs_created > refresh_browser_tabs)
    ):
        if context:
            await context.close()
        if browser:
            await browser.close()
        browser = await get_browser_async(
            playwright, executable_path, headless=headless, incognito=incognito
        )
        context = await browser.new_context(
            ignore_https_errors=True, bypass_csp=True, accept_downloads=False
        )

    page = await context.new_page()
    viewport = {
        "width": width or page.viewport_size["width"],
        "height": height or page.viewport_size["height"],
    }
    await page.set_viewport_size(viewport_size=viewport)
    return page

async def close_playwright():
    """
    Closes the browser context, browser, and stops Playwright.
    """
    await context.close() # Close the browser context
    await browser.close() # Close the browser
    await playwright.stop() # Stop Playwright

return new_page, close_playwright

```

Sqlite Database Manager

```
class Status(Enum):
    """Enumeration for possible link statuses in the database."""

    DOWNLOADED = "downloaded"
    CRAWL = "crawl"
    CRAWLED = "crawled"

class DatabaseManager:
    """
    Manages a SQLite database of links with status, count, and timestamps.
    """

    def __init__(self, db_name: str):
        """Initialize database and Link model."""
        self.db = peewee.SqliteDatabase(db_name)

        class Link(peewee.Model):
            link = peewee.CharField(primary_key=True)
            status = peewee.CharField(choices=[s.value for s in Status])
            count = peewee.IntegerField(default=0)
            created_at = peewee.DateTimeField(
                default=lambda: datetime.now(timezone.utc)
            )
            updated_at = peewee.DateTimeField(
                default=lambda: datetime.now(timezone.utc)
            )

            class Meta:
                database = self.db

            def save(self, *args, **kwargs):
                self.updated_at = datetime.now(timezone.utc)
                return super().save(*args, **kwargs)

        self.Link = Link
        self.db.create_tables([self.Link])

    def reload_connection(self) -> bool:
        """Reload the database connection to recover from issues."""
```

```

try:
    self.close()
    self.__init__(self.db.database)
    return True
except Exception as e:
    print(f"Error reloading database connection: {e}")
    return False

def create(self, link: str, status: Status, count: int = 0) -> bool:
    """Insert a new link with status and count."""
    try:
        self.Link.create(link=link, status=status.value, count=count)
        return True
    except peewee.IntegrityError:
        return False

def read(
    self,
    link: str = None,
    status: Status = None,
    count: int = None,
    limit: int = None,
    count_expression: Literal["gt", "lt", "eq"] = "gt",
    include_timestamps: bool = False,
    randomize: bool = False,
) -> List[Tuple]:
    """
    Read links with optional filters and limit.
    Returns tuples of (link, status, count[, created_at, updated_at]).
    """
    query = self.Link.select()
    if link:
        query = query.where(self.Link.link == link)
    if status:
        query = query.where(self.Link.status == status.value)
    if count is not None:
        if count_expression == "gt":
            query = query.where(self.Link.count > count)
        elif count_expression == "lt":
            query = query.where(self.Link.count < count)
        else:
            query = query.where(self.Link.count == count)

```

```

    if randomize:
        query = query.order_by(peewee.fn.Random())
    if limit:
        query = query.limit(limit)
    if include_timestamps:
        return [
            (
                l.link,
                l.status,
                l.count,
                l.created_at.isoformat(),
                l.updated_at.isoformat(),
            )
            for l in query
        ]
    else:
        return [(l.link, l.status, l.count) for l in query]

def read_where(
    self,
    limit: int = None,
    include_timestamps: bool = False,
    randomize: bool = False,
    **filters,
) -> List[Tuple]:
    """
    Flexible read with keyword filters.
    Returns tuples of (link, status, count[, created_at, updated_at]).
    """
    query = self.Link.select()
    for field, value in filters.items():
        if field == "link":
            query = query.where(self.Link.link == value)
        elif field == "status":
            if isinstance(value, Status):
                value = value.value
            query = query.where(self.Link.status == value)
        elif field == "link_contains":
            query = query.where(self.Link.link.contains(value))
        elif field == "link_startswith":
            query = query.where(self.Link.link.startswith(value))
        elif field == "count":

```

```

        query = query.where(self.Link.count == value)
    elif field == "count_gt":
        query = query.where(self.Link.count > value)
    elif field == "created_after":
        if isinstance(value, str):
            value = datetime.fromisoformat(value.replace("Z", "+00:00"))
        query = query.where(self.Link.created_at > value)
    elif field == "created_before":
        if isinstance(value, str):
            value = datetime.fromisoformat(value.replace("Z", "+00:00"))
        query = query.where(self.Link.created_at < value)
    if randomize:
        query = query.order_by(peewee.fn.Random())
    if limit:
        query = query.limit(limit)
    if include_timestamps:
        return [
            (
                l.link,
                l.status,
                l.count,
                l.created_at.isoformat(),
                l.updated_at.isoformat(),
            )
            for l in query
        ]
    else:
        return [(l.link, l.status, l.count) for l in query]

def update(self, link: str, status: Status = None, count: int = None) -> bool:
    """Update status and/or count for a link. Returns True if updated."""
    updates = {}
    if status:
        updates["status"] = status.value
    if count is not None:
        updates["count"] = count
    updates["updated_at"] = datetime.now(timezone.utc)
    if len(updates) == 1:
        return False
    query = self.Link.update(**updates).where(self.Link.link == link)
    return query.execute() > 0

```

```

def increment_count(self, link: str) -> bool:
    """Increment the count for a link. Returns True if updated."""
    query = self.Link.update(
        count=self.Link.count + 1, updated_at=datetime.now(timezone.utc)
    ).where(self.Link.link == link)
    return query.execute() > 0

def delete(self, link: str) -> bool:
    """Delete a link. Returns True if deleted."""
    return self.Link.delete().where(self.Link.link == link).execute() > 0

def get_statistics(self) -> dict:
    """Return statistics: total, per-status, oldest/newest timestamps."""
    total_count = self.Link.select().count()
    status_counts = {}
    for status in Status:
        count = self.Link.select().where(self.Link.status == status.value).count()
        status_counts[status.value] = count
    oldest = self.Link.select().order_by(self.Link.created_at.asc()).first()
    newest = self.Link.select().order_by(self.Link.created_at.desc()).first()
    return {
        "total_links": total_count,
        "status_breakdown": status_counts,
        "oldest_entry": oldest.created_at if oldest else None,
        "newest_entry": newest.created_at if newest else None,
    }

def close(self):
    """Close the database connection."""
    if not self.db.is_closed():
        self.db.close()

```

Schema

```

class PagePattern(TypedDict):
    """
    Represents a pattern for matching or ignoring pages.
    - regex: Regular expression to match URLs.
    - examples: Example URLs matching the regex.
    """

```

```

    regex: str
    examples: List[str]

class Schema(TypedDict):
    """
    Defines crawling rules for a website.
    - prepare_page_expression: JavaScript code to run before extracting links.
    - locations: List of allowed page patterns.
    - ignorePages: List of patterns to ignore.
    """
    prepare_page_expression: str
    locations: List[PagePattern]
    ignorePages: List[PagePattern]

schemas: List[Schema] = [
    {
        "prepare_page_expression": "() => Promise.resolve()",
        "locations": [
            {"regex": ".*kenyalaw.org.*", "examples": ["https://kenyalaw.org/"]}
        ],
        "ignorePages": [
            {
                "regex": "^(?!.*kenyalaw\\.org).*",
                "examples": ["https://www.facebook.com/MyKenyaLaw"],
            },
        ],
    }
]

```

The following function serializes Python data structures into a string suitable for JavaScript, treating string values of specified **raw_keys** as raw code (not quoted).

```

def custom_serialize(data: Any, raw_keys: list[str] = None):
    """
    Recursively serializes Python data structures into a string
    suitable for JavaScript, treating string values of specified 'raw_keys'
    as raw code (not quoted).

    Args:
        data: The Python data structure (dict, list, str, int, float, bool, None)
              to be serialized.
        raw_keys: An optional list of strings representing keys whose string values
    """

```

```

        should be treated as raw JavaScript code and not quoted.
"""
if raw_keys is None:
    raw_keys = []

if isinstance(data, dict):
    # If it's a dictionary, serialize its keys and values
    items = []
    for k, v in data.items():
        # Keys are usually strings and should be quoted if not raw
        key_str = (
            json.dumps(k) if isinstance(k, str) else str(k)
        ) # Use json.dumps for keys too for consistency

        # Check if the current key is in the raw_keys list and its value is a string
        if k in raw_keys and isinstance(v, str):
            # If it's a raw key and a string, return the value directly without quoting
            items.append(f"{key_str}: {v}")
        else:
            # Otherwise, recursively serialize the value with the same raw_keys list
            items.append(f"{key_str}: {custom_serialize(v, raw_keys)}")
    return "[" + ", ".join(items) + "]"
elif isinstance(data, list):
    # If it's a list, serialize each item recursively
    elements = [custom_serialize(item, raw_keys) for item in data]
    return "[" + ", ".join(elements) + "]"
elif isinstance(data, str):
    # Regular strings need to be quoted using json.dumps to handle escaping
    return json.dumps(data)
elif isinstance(data, (int, float, bool, type(None))):
    # Numbers, booleans, and None (null in JS) can be converted directly
    return (
        str(data).lower() if isinstance(data, bool) else str(data)
    ) # Convert Python True/False to JS true/false
else:
    # Fallback for other unsupported types
    raise TypeError(f"Unsupported type for serialization: {type(data)}")

```

Generate the JavaScript code to evaluate the schemas and extract links:

```

def generate_evaluate_script(schemas: List[Schema]):
    schemas_str = custom_serialize(schemas, ["prepare_page_expression"])
    evaluate_script = (
        "async () => {\n"
        f"const schemas = {schemas_str}\n"
        """"
        /** @type {string[]} */
        const errors = [];
        const generator = async function* () {
            try {
                for (const schema of schemas) {
                    try {
                        const isLocationAllowed = schema.locations.some((domain) =>
                            new RegExp(domain.regex).test(location.host + location.pathname)
                        );
                        if (isLocationAllowed) {
                            if (schema.prepare_page_expression) {
                                await Promise.resolve(schema.prepare_page_expression());
                            }

                            const $anchors = [...document.querySelectorAll("a")];
                            /**
                             * Validates string
                             * @param {HTMLAnchorElement} link
                             * @returns {string} the valid string
                             */
                            const validateLink = (link) => {
                                try {
                                    if (link.href) {
                                        const url = new URL(link.href);
                                        const validProtocols = ["https:", "http:", "ftp:"];
                                        const normalizedProtocol = url.protocol.toLowerCase();
                                        if (validProtocols.includes(normalizedProtocol)) {
                                            return link.href;
                                        }
                                    }
                                } catch (error) {
                                    errors.push(error);
                                }
                                return "";
                            };
                            /**

```

```

    * @param {string} link
    * @returns {boolean}
    */
    const isIgnoreLink = (link) => {
        return schema.ignorePages.some((ignorePage) => {
            return new RegExp(ignorePage.regex).test(new URL(link).host);
        });
    };
    const links = [
        ...new Set($anchors.map(validateLink).filter((i) => !!i)),
    ];

    for (const link of links) {
        try {
            if (!isIgnoreLink(link)) {
                yield link;
            }
        } catch (error) {
            errors.push(error);
        }
    }
} catch (error) {
    errors.push(error);
}
} catch (error) {
    errors.push(error);
}
};

/** @type {string[]} */
const links = [];
for await (const item of generator()) {
    links.push(item);
}
return [links, errors];
}
"""
)
return evaluate_script

```

```
print(jsbeautifier.beautify(generate_evaluate_script(schemas)))
```

```
async () => {
  const schemas = [{
    "prepare_page_expression": () => Promise.resolve(),
    "locations": [{
      "regex": ".*kenyalaw.org.*",
      "examples": ["https://kenyalaw.org/"]
    }],
    "ignorePages": [{
      "regex": "^(?!.*kenyalaw\\.org).*",
      "examples": ["https://www.facebook.com/MyKenyaLaw"]
    }]
  }]

  /** @type {string[]} */
  const errors = [];
  const generator = async function*() {
    try {
      for (const schema of schemas) {
        try {
          const isLocationAllowed = schema.locations.some((domain) =>
            new RegExp(domain.regex).test(location.host + location.pathname)
          );
          if (isLocationAllowed) {
            if (schema.prepare_page_expression) {
              await Promise.resolve(schema.prepare_page_expression());
            }

            const $anchors = [...document.querySelectorAll("a")];
            /**
             * Validates string
             * @param {HTMLAnchorElement} link
             * @returns {string} the valid string
             */
            const validateLink = (link) => {
              try {
                if (link.href) {
                  const url = new URL(link.href);
                  const validProtocols = ["https:", "http:", "ftp:"];
                  const normalizedProtocol = url.protocol.toLowerCase();
                  if (validProtocols.includes(normalizedProtocol)) {
```

```

        return link.href;
    }
}
} catch (error) {
    errors.push(error);
}
return "";
};
/**
 * @param {string} link
 * @returns {boolean}
 */
const isIgnoreLink = (link) => {
    return schema.ignorePages.some((ignorePage) => {
        return new RegExp(ignorePage.regex).test(new URL(link).host)
    });
};
const links = [
    ...new Set($anchors.map(validateLink).filter((i) => !!i)),
];

for (const link of links) {
    try {
        if (!isIgnoreLink(link)) {
            yield link;
        }
    } catch (error) {
        errors.push(error);
    }
}
} catch (error) {
    errors.push(error);
}
} catch (error) {
    errors.push(error);
}
};

/** @type {string[]} */
const links = [];
for await (const item of generator()) {

```

```

        links.push(item);
    }
    return [links, errors];
}

```

Initialize the Crawler

```

# Ensure the database path exists
database_path = Path(database_path_str)
database_path.parent.mkdir(exist_ok=True)
# Initialize the database manager
database = DatabaseManager(database_path)

```

```

# Initialize Playwright page factory and cleanup function
new_page, close_playwright = await get_browser_page_async(
    refresh_browser_tabs=refresh_browser_tabs
)

```

```

async def crawl_and_save(link: str):

```

```

    """

```

```

    Crawls a single URL, extracts sub-links using the schema-based JS,
    saves new links to the database, and updates crawl status.
    Retries on error by incrementing the count.
    """

```

```

    try:

```

```

        page: Page = await new_page()
        await page.goto(link, wait_until="networkidle")

```

```

        sub_links, errors = await page.evaluate(generate_evaluate_script(schemas))

```

```

        if len(errors) > 0:

```

```

            print(errors)

```

```

        await page.close()

```

```

        added_count = 0

```

```

        for sub_link in sub_links:

```

```

            if database.create(sub_link, Status.CRAWL):

```

```

                added_count = added_count + 1

```

```

        database.update(link, Status.CRAWLED)

```

```

        if added_count > 0:

```

```

            print(f"Added {added_count}/{len(sub_links)} links from {link}")

```

```

    except Exception as e:

```

```
print(e)
database.increment_count(link)
```

```
async def crawl_and_download():
    """
    Recursively crawls up to `crawl_count` URLs with status=Status.CRAWL and count=0.
    For each batch:
        - Runs `crawl_and_save` concurrently on all selected records.
        - Reloads the database connection after each batch.
        - Continues until no more URLs to crawl.
    """
    records = database.read(
        limit=crawl_count,
        status=Status.CRAWL,
        count=0,
        count_expression="eq",
        randomize=True)
    if len(records) <= 0:
        print("Stopped...")
        return
    awaitables = [crawl_and_save(record[0]) for record in records]
    await asyncio.gather(*awaitables)
    await asyncio.sleep(1)
    database.reload_connection()
    await crawl_and_download()
```

```
# # Initialize the crawling process
# await crawl_and_save("https://kenyalaw.org/caselaw/cases/advanced_search/")
```

Added 323/323 links from https://kenyalaw.org/caselaw/cases/advanced_search/

```
# Start the crawling process
await crawl_and_download()
```

Added 59/68 links from <https://www.kenyalaw.org/kl/index.php?id=3158>

Added 23/91 links from https://kenyalaw.org/caselaw/cases/advanced_search/?totaljudge=231&cor

Added 19/83 links from https://kenyalaw.org/caselaw/cases/advanced_search/?totaljudge=249&cor

```
# Print database statistics
database.get_statistics()
```

```
{'total_links': 424,
 'status_breakdown': {'downloaded': 0, 'crawl': 419, 'crawled': 5},
 'oldest_entry': '2025-08-04 15:57:56.569805+00:00',
 'newest_entry': '2025-08-04 15:58:15.964191+00:00'}
```

Conclusion

The beauty of this system lies in its adaptability. Start with a simple schema for your target site, define what to crawl and what to ignore, and let the system discover the structure for you.

Whether you're building a legal research tool, academic database, or specialized content aggregator, schema-based crawling gives you the surgical precision that enterprise tools can't match, without the enterprise complexity.

After crawling, you can examine which links to process further. For example, you can filter links that match cases or legislation such as <http://kenyalaw.org/caselaw/cases/view/187364> and decide whether to download and process the PDF (<http://kenyalaw.org/caselaw/cases/export/187364/pdf>) or Word document (<http://kenyalaw.org/caselaw/cases/export/187364/doc>) or XML (<http://kenyalaw.org/caselaw/cases/export/187364/xml>).

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)