


CodeGraph: Turn Your Codebase into a Knowledge Graph So AI Agents Stop Wasting Tokens

Kabui, Charles

2026-06-03

[Read at ToKnow.ai](#)



The banner features a dark blue background with a subtle grid pattern. On the right side, there is a faint, stylized graph icon with three nodes and connecting lines. The main text is in white and light blue. Below the title, there are three statistics presented in colored boxes with vertical bars on the left: a red box for 47%, a yellow box for 58%, and a light blue box for 20+.

CodeGraph: Turn Your Codebase into a Knowledge Graph
Pre-indexed code intelligence for AI coding agents

- 47%**
Fewer tokens used across 7 codebases
- 58%**
Fewer tool calls, 81% on VS Code repo
- 20+**
Languages parsed via tree-sitter, 100% local

June 3, 2026

ToKnow.ai

Every time you start a new AI coding session, the agent re-reads your project from scratch, burning tokens on file discovery. [CodeGraph](#) fixes this by pre-indexing your repository into a

local SQLite-backed knowledge graph that agents query through an [MCP server](#). It uses tree-sitter to parse 20+ languages, extracts symbol relationships, call graphs, and dependency edges, then serves them on demand. Benchmarked across seven open-source codebases including VS Code (~10k files) and Django (~3k files), CodeGraph cut token usage by 47%, tool calls by 58%, and wall-clock time by 22% on average. On VS Code specifically, tool calls dropped 81%. It plugs into Claude Code, Cursor, Codex, Gemini CLI, OpenCode, and Kiro with a single `codegraph install` command. A file watcher keeps the graph current as you code, and everything runs locally with no external API calls.

Tokens cost money. At Opus 4.8 pricing (\$5/\$25 per million input/output tokens), a developer exploring a large codebase burns real dollars on context loading. CodeGraph replaces those exploratory file reads with instant graph queries: a structural question like “what depends on this module?” that would require reading 20+ files gets answered in a fraction of the tokens. It also handles cross-language boundaries that trip up static analysis, bridging Swift/Objective-C calls, React Native bridges, TurboModules, and Expo modules.

AI coding agents are maturing past raw model improvements into tooling that makes each token more productive. [Dynamic workflows](#) scale agents horizontally. CodeGraph makes each individual agent session cheaper and faster by front-loading structural understanding.

Sources:

- [CodeGraph GitHub Repository](#)
- [CodeGraph Documentation](#)
- [CodeGraph npm Package](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)