

Repeated Random Sampling

Approximating Distributions & Estimating Pi Using Monte Carlo Techniques in Python

Kabui, Charles

2025-05-15

Table of Contents

Examples	4
Estimating with Random Sampling	4
Solution	5
Estimating the Mean of a Normal Distribution	6
Solution	6
Estimating Dice Probability	8
Solution	8

[!\[\]\(e3f8612927870f2e0f9f5989e6dd3064_img.jpg\) Read at ToKnow.ai](#)[!\[\]\(003082e50e3009141f59bd5df831749f_img.jpg\) Open in Kaggle](#)[!\[\]\(17413706fd4997a1a4bdf85c6864eee1_img.jpg\) Download as Notebook](#)

Random Sampling is a fundamental statistical technique for selecting a subset of data from a larger population. Each sample is chosen randomly and independently. It's crucial in statistics and data science for providing unbiased insights.

Repeated Random Sampling is emphasized because a single sample may not represent the population well. By repeating the sampling process many times, variability is reduced, improving the approximation of true population parameters and leading to more robust estimates. This process forms the basis for bootstrapping and **Monte Carlo simulations**.

Monte Carlo methods are computational algorithms that use random sampling to obtain numerical results, often useful when deterministic methods are too complex. They rely on

generating random numbers as inputs, running simulations multiple times, representing uncertainty with probability distributions, and analyzing outputs.

A key application is **Approximating Distributions**. Thousands of random samples can be simulated from a distribution to understand its behavior, visualize histograms, calculate statistics, and construct confidence intervals. This involves using pseudo-random number generation to transform uniformly distributed numbers into desired distributions.

Repeated sampling helps identify patterns and converge to population characteristics. These techniques are also relevant to **Hypothesis Testing And Estimation**.

A classic Monte Carlo application is **Estimating Pi**. This involves simulating random points within a unit square with an inscribed quarter circle. The ratio of points falling inside the quarter circle (where $x^2 + y^2 \leq 1$) to the total number of points approximates $\frac{\pi}{4}$. As the number of points (N) increases, the estimate becomes more accurate. Visualizing these points shows the quarter circle becoming more defined with more samples.

Error Analysis and Convergence show that the error between the estimated and true Pi decreases as N increases. This convergence follows a law where the Error is proportional to $\frac{1}{\sqrt{N}}$. This indicates that while increasing N improves precision, there are diminishing returns.

```
import random
import matplotlib.pyplot as plt

N = 10000
x, y = (
    [random.random() for _ in range(N)],
    [random.random() for _ in range(N)])
inside = [
    (x[i], y[i])
    for i
    in range(N) if x[i]**2 + y[i]**2 <= 1]
pi = 4 * len(inside) / N
print(f"Estimated Pi: {pi:.5f}")


plt.figure(figsize=(6,6))
plt.scatter(
    *zip(*inside),
    color='blue',
    s=1,
    label='inside')
plt.scatter(
    *zip(*[(x[i], y[i]) for i in range(N) if (x[i], y[i]) not in inside]),
    color='red',
```

```

    s=1,
    label='outside')
plt.legend(
    loc = "upper right",
    facecolor="green",
    edgecolor='white',
    markerscale=14,
    labelcolor='white',
    prop={'weight': 'bold', 'size': 14})
plt.title(f"Monte Carlo Pi Estimate\nEstimated Pi   {pi:.5f}")
plt.xlabel("x")
plt.ylabel("y")
plt.axis("equal")
plt.show()

```

Estimated Pi: 3.11560



Beyond estimating Pi, **Monte Carlo and random sampling** methods are used in various domains, including numerical integration, financial risk modeling, physics simulations, and importantly for data science, **predictive analytics and machine learning** (e.g., bagging).

Examples

Estimating π with Random Sampling

You are designing a simulation to estimate the value of π using 100,000 random points uniformly sampled inside a unit square.


How accurate is your estimate compared to the true value of π ?
Use random sampling and compare the estimate over 10 different runs.

Solution

```
import numpy as np
import matplotlib.pyplot as plt
from IPython.display import Markdown

estimates = []
for _ in range(10):
    x, y = np.random.rand(100000), np.random.rand(100000)
    pi_estimate = 4 * np.mean(x**2 + y**2 <= 1)
    estimates.append(pi_estimate)

Markdown(f"Estimates of  $\pi$  from 10 runs:\n{np.array(estimates).tolist()}")
# Plotting the estimates
plt.plot(estimates, marker='o', label='Estimated  $\pi$ ')
plt.axhline(np.pi, color='r', linestyle='--', label=' $\pi$ ', alpha=0.7)
plt.title('Estimates of  $\pi$  over 10 Runs')
plt.ylabel('Estimated  $\pi$ ')
plt.xlabel('Run')
plt.legend()
plt.show()
```



The estimate of π is around the true value of π (3.14159) with a margin of error.

Estimating the Mean of a Normal Distribution

Suppose you repeatedly draw 1,000 samples of size 50 from a normal distribution with mean $\mu = 10$ and standard deviation $\sigma = 2$.

What does the distribution of the sample means look like?

Plot a histogram of the sample means and compute the overall average and standard deviation.

Solution


```
means = [np.mean(np.random.normal(10, 2, 50)) for _ in range(1000)]

print(f"Average of means: {np.mean(means):.4f}")
print(f"Std dev of means: {np.std(means):.4f}")
```

```
# Plotting the histogram of sample means
plt.hist(means, bins=30, edgecolor='black', align='left', alpha=0.7)
plt.title('Histogram of Sample Means')
plt.xlabel('Mean')
plt.ylabel('Frequency')
plt.show()
```

Average of means: 10.0075

Std dev of means: 0.2896



Average of means ≈ 10

This aligns well with the true population mean ($\mu = 10$), as expected from the **Law of Large Numbers**.

Standard deviation of means ≈ 0.2896

This closely matches the **theoretical standard error**:

$$SE = \frac{\sigma}{\sqrt{n}} = \frac{2}{\sqrt{50}} \approx 0.2828$$

This shows that your sampling distribution of the mean is behaving exactly as it should.

Estimating Dice Probability

Two fair six-sided dice are rolled.

Using 100,000 simulations, estimate the probability that the sum of the two dice is greater than 9.

How close is this to the theoretical probability?

Solution

Theoretical probability:

There are 6 outcomes > 9 out of 36 possible outcomes

When rolling two dice, the sum can be greater than 9 in these cases:

- **Sum of 10:** (4,6), (5,5), (6,4)
- **Sum of 11:** (5,6), (6,5)
- **Sum of 12:** (6,6)

$$P(\text{sum} > 9) = \frac{6}{36} = 0.1667$$

```
sums = np.random.randint(1, 7, (100000, 2)).sum(axis=1)
prob = np.mean(sums > 9)
print(f"Estimated P(sum > 9): {prob:.4f}")
print(f"Theoretical P(sum > 9): {6/36:.4f}")
```

Estimated $P(\text{sum} > 9)$: 0.1661

Theoretical $P(\text{sum} > 9)$: 0.1667

The estimated probability from the simulation is close to the theoretical probability, with a small margin of error, due to **random sampling error**.

***Disclaimer:** For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)*