

daVinci-MagiHuman: One Model Generates Synchronized Video and Audio in 2 Seconds

Kabui, Charles

2026-03-29

[Read at ToKnow.ai](#)



SII-GAIR and Sand.ai released [daVinci-MagiHuman](#), a 15B-parameter open-source model that jointly generates synchronized video and audio from a single text prompt. Most audio-video systems use complex multi-stream architectures with cross-attention to keep different outputs in sync. daVinci-MagiHuman uses a single-stream Transformer instead: text, video, and audio are all processed as one unified token sequence using only self-attention. A “sandwich” layout

gives the first and last four layers modality-specific projections while the middle 32 layers share parameters across all modalities. Combined with model distillation (down to 8 denoising steps), latent-space super-resolution, and a Turbo VAE decoder, the system generates a 5-second 256p video in [2 seconds on a single H100 GPU](#). In human evaluation across 2,000 pairwise comparisons, it achieved an 80.0% win rate against Ovi 1.1 and 60.9% against LTX 2.3, with the lowest word error rate (14.60%) for speech intelligibility among leading open models. It supports six languages: English, Mandarin, Cantonese, Japanese, Korean, German, and French.

That 2-second generation time matters. Near-real-time audio-video generation means content creators and developers can iterate on talking-head videos, multilingual voiceovers, or virtual assistant prototypes without waiting minutes per clip. The full stack is open-sourced under Apache 2.0, including the base model, distilled model, super-resolution model, and inference code, so anyone with an H100 (or even [consumer GPUs with some adjustments](#)) can run it.

The broader pattern here is striking: simpler architectures keep winning. Cross-attention, multi-stream coordination, and modality-specific pipelines are giving way to unified single-stream designs. ByteDance’s [Seedance](#) took a similar multimodal approach, but daVinci-MagiHuman pushes architectural simplicity further while matching or exceeding quality.

Sources:

- [daVinci-MagiHuman Paper \(arXiv\)](#)
- [daVinci-MagiHuman GitHub \(1.1k stars\)](#)
- [HuggingFace Model Page](#)
- [HuggingFace Paper Discussion \(115 upvotes\)](#)
- [HuggingFace Demo Space](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)