

Gemini 3.7 Flash: Google's Workhorse Coding Model at Half the Price, Three Weeks After 3.6

Kabui, Charles

2026-08-20

[Read at ToKnow.ai](#)



The graphic features a dark blue background with a network diagram of three nodes in the top right. The main title is in large white and teal font. Below it, a subtitle in white text states 'Three weeks after 3.6 Flash, with better coding and agent skills'. Three key metrics are highlighted in separate boxes with vertical bars: a 43.6% increase in coding score, a price of \$0.75 per million tokens, and a 3-week timeline. The date 'August 20, 2026' and the source 'ToKnow.ai' are at the bottom.

Gemini 3.7 Flash: Google's Workhorse Model at Half the Price

Three weeks after 3.6 Flash, with better coding and agent skills

43.6% FrontierCode 1.1 coding score, up from 34.4%	\$0.75 per million input tokens, half of 3.6 Flash	3 weeks after Gemini 3.6 Flash, with a 1M context window
---	---	---

August 20, 2026 ToKnow.ai

Google released [Gemini 3.7 Flash](#) on August 13, its most capable workhorse model for coding and agents, three weeks after Gemini 3.6 Flash. A workhorse model is the fast, cheaper option for high-volume everyday work, while bigger models handle the hardest problems. 3.7

Flash is built on 3.6 Flash with stronger reasoning, and ships at an introductory price of \$0.75 per million input tokens and \$3.75 per million output tokens, half of 3.6 Flash. On [FrontierCode 1.1](#) it scores 43.6% versus 34.4%, and on [DeepSWE v1.1](#) it reaches 65.3% versus 49.0%. Document reasoning jumps on GDP.pdf (34.0% versus 22.0%) and business workflow automation on AutomationBench (30.4% versus 17.0%). It keeps a 1 million token context window.

The practical win is price. A team running agents on 3.6 Flash gets better coding and tool use on 3.7 Flash at roughly half the input cost, important when agents send many small requests. The introductory rate runs through December 31, 2026, then rises to \$1.50 per million input tokens, so the savings are time-limited. Google is also moving [Gemini Spark](#), its 24/7 personal agent for AI Pro and Ultra subscribers, onto 3.7 Flash, making it more accurate on multi-step tasks in Google Workspace apps.

The release is also a statement about pace. Google has shipped a meaningfully better workhorse model three weeks after the previous one, and the biggest gains landed in the cheap tier, not the flagship. For everyday agent work, the frontier is moving fastest at the bottom of the price curve.

Read More: [Gemini 3.1 Flash-Lite: Google's Fastest Model Costs \\$0.25 per Million Tokens](#)

Sources:

- [Introducing Gemini 3.7 Flash \(Google\)](#)
- [Gemini 3.7 Flash model card \(Google DeepMind\)](#)
- [Gemini Developer API pricing \(Google\)](#)
- [FrontierCode benchmark \(Cognition\)](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)