Gemini 4 Argon: A Million-Token Model That Ships to Cyber Defenders First

Kabui, Charles

2026-10-08

Read at ToKnow.ai



Google released Gemini 4 Argon on September 30, the first model in its Gemini 4 line and its answer to OpenAI's GPT-6 Astra. It is built for long, multi-step work such as software engineering, legal and financial research, and cybersecurity defense. The biggest change is its output limit, raised to a million tokens from 64,000, which lets the model write hundreds of thousands of tokens in one pass. Google reports 77.9% on DeepSWE v1.1, a long-horizon

coding test, and 91.7% on LVBench for long video. Introductory pricing is $2 per million input tokens and $10 per million output, rising later to $4 and $20.

Most people cannot use it yet. Google is giving Argon first to vetted cyber defenders through its Fairwind Program, which it says has over 650 partners, and to its own engineers. Security firm Wiz reports the model found a critical flaw in hospital software that earlier frontier models missed. Two outside checks agree it is frontier-class: Artificial Analysis scores it 53 on its Intelligence Index, level with GPT-6 Astra, and it ties for first on the CWE-bench security benchmark at 68%.

Argon reasons in one long pass rather than exploring several reasoning paths at once, the approach some researchers call parallel thinking. Google is also shipping it in an unusual order: defenders and governments get Argon first, while everyone else waits for a date the company has not named, beyond "as soon as possible."

Read More: Google's earlier defender-only model

Sources:

- Introducing Gemini 4 Argon (Google)
- Fairwind Program (Google DeepMind)
- CWE-bench v1 leaderboard (Collinear AI)
- Gemini 4 Argon (Artificial Analysis)
- Gemini 4 Argon is Google's most ambitious AI model yet (Digital Trends)

Disclaimer: *For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation.* ***Read more:*** */terms-of-service*