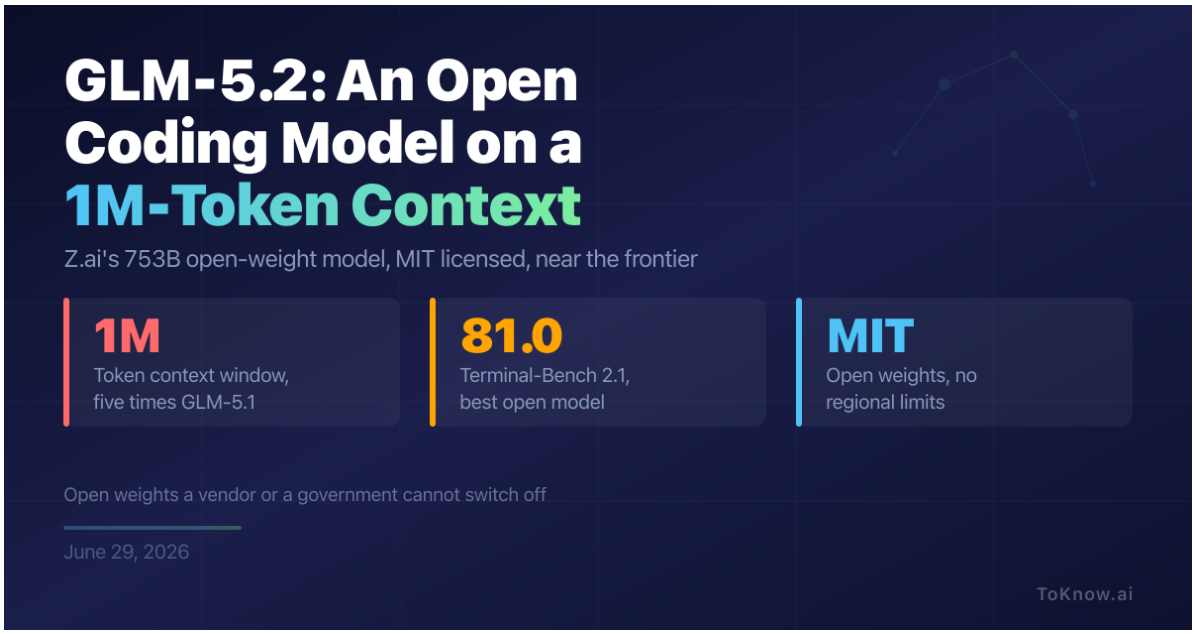
GLM-5.2: An Open-Weight Coding Model With a 1M-Token Context That Rivals the Frontier

Kabui, Charles

2026-06-29

Read at ToKnow.ai



Z.ai released GLM-5.2, an open-weight model built for coding jobs that run for hours. It is a 753-billion-parameter mixture-of-experts model, meaning it holds many specialized sub-networks but activates only a small slice per token to stay fast, and it ships under the MIT

license with no regional restrictions. The biggest change is a working 1M-token context, five times the 200K of GLM-5.1 and enough to hold a whole mid-size codebase at once. A method called IndexShare cuts the compute per token by 2.9 times at that length, so the long context stays cheap to run. On FrontierSWE, a test of open-ended projects that take hours, it scores 74.4, about 1% behind Claude Opus 4.8 and just ahead of GPT-5.5. On Terminal-Bench 2.1 it jumps to 81.0 from GLM-5.1's 63.5, the best open-source result and near Opus 4.8's 85.0.

Because the weights are downloadable and MIT-licensed, a team can self-host a model that competes with the best closed coding agents on multi-hour work, with no per-token API bill and no country-based lockout. The large context helps agents most: instead of re-reading files repeatedly, the model can hold an entire project and its history in one pass.

It landed the same month a White House order briefly switched off access to Anthropic's Mythos and Fable 5 models in South Korea. An open-weight model you can download cannot be revoked by a vendor or a government, which makes the open versus closed gap increasingly about ownership and resilience, not raw capability.

Read More: GLM-5, the predecessor that first closed the open-source gap

Sources:

- GLM-5.2: Built for Long-Horizon Tasks (Z.ai)
- GLM-5.2 model weights on Hugging Face
- GLM-5 GitHub repository
- GLM-5.2 model guide (Z.ai docs)

Disclaimer: *For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation.* ***Read more:*** */terms-of-service*