

Google DeepMind Lyria 3: Music Generation Model Outputs 48kHz Stereo with Real-Time Steering and SynthID Watermarking

Kabui, Charles

2026-03-03

[Read at ToKnow.ai](#)



The image is a promotional graphic for Google DeepMind Lyria 3. It features a dark blue background with a grid pattern. At the top left, the text "Google DeepMind Lyria 3" is displayed in white and green. Below this, a subtitle reads "AI music generation with real-time steering and SynthID watermarking". The graphic is divided into three main sections by vertical lines. The first section, on the left, is titled "48kHz" in red and lists "Studio-quality stereo" and "Full arrangements from scratch". The second section, in the middle, is titled "<2s" in yellow and lists "Real-time latency" and "Steer music as it plays". The third section, on the right, is titled "SynthID" in blue and lists "Inaudible watermark" and "Survives compression". At the bottom left, a line of text states "Text, image, or audio → full multi-instrumental tracks with vocals". In the bottom right corner, the "ToKnow.ai" logo is visible. A faint, stylized graphic of a person's head and shoulders is in the top right corner.

Google DeepMind
Lyria 3

AI music generation with real-time steering and SynthID watermarking

48kHz
Studio-quality stereo
Full arrangements from scratch

<2s
Real-time latency
Steer music as it plays

SynthID
Inaudible watermark
Survives compression

Text, image, or audio → full multi-instrumental tracks with vocals

ToKnow.ai

Figure 1: Google DeepMind Lyria 3

Google DeepMind released Lyria 3, its most advanced music generation model. The model generates 30-second tracks in studio-quality stereo (48kHz) from text, image, or audio inputs. It does not piece together loops. It generates full musical arrangements from scratch, including vocals and multi-instrumental tracks. Lyria 3 is now integrated into the Gemini app. Describe a mood, genre, or set of instruments, and it outputs a finished audio file. Music generation is hard because it is continuous and multi-layered: melody, harmony, rhythm, and timbre all need to stay coherent from the first second to the last.

The real shift is the Lyria RealTime API. Traditional models work like a jukebox: input a prompt, wait for a file. Lyria RealTime generates audio in 2-second chunks as a live stream. Users can steer style, tempo, and instrumentation while the track plays, with under 2 seconds of delay. The Music AI Sandbox builds on this, transforming a hum into a full orchestral arrangement or turning chord progressions into a vocal choir. Every track includes a SynthID watermark embedded in the audio waveform, inaudible to humans but detectable by software even after MP3 compression or analog re-recording. As AI-generated music floods streaming platforms, permanent attribution solves a real trust problem for creators and listeners. Lyria 3 competes directly with Suno and Udio but brings multimodal input and Google's infrastructure advantage.

Sources:

- [Google DeepMind Lyria Technical Details](#)
- [MarkTechPost: Google DeepMind Lyria 3](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)