

Google Jules: The Coding Agent That Writes Pull Requests in the Background

Kabui, Charles

2026-05-28

[Read at ToKnow.ai](#)

The screenshot shows the Google Jules website with a dark blue background. The main heading is "Google Jules: The Coding Agent That Writes Pull Requests for You", with "Pull Requests for You" in a light blue color. Below the heading, it says "Asynchronous, cloud VM, Gemini 3.1 Pro". There are three feature boxes: "300/day Tasks per day on the Jules Ultra tier", "60 Concurrent agent tasks running in parallel", and "-9.5% Task failures after the Planning Critic agent". At the bottom left, it says "May 28, 2026" and at the bottom right, "ToKnow.ai".

Google Jules: The Coding Agent That Writes Pull Requests for You

Asynchronous, cloud VM, Gemini 3.1 Pro

- 300/day**
Tasks per day on the Jules Ultra tier
- 60**
Concurrent agent tasks running in parallel
- 9.5%**
Task failures after the Planning Critic agent

May 28, 2026

ToKnow.ai

[Jules](#) is Google’s autonomous coding agent, powered by [Gemini 3.1 Pro](#). Unlike Cursor or Claude Code, which sit in your editor waiting for the next prompt, Jules runs asynchronously. You give it a task such as bumping Next.js to 15, and Jules clones your repo into a Google Cloud VM, drafts a plan you can edit, writes the code, and opens a pull request on GitHub. A “Planning Critic” agent reviews every auto-approved plan, which Google says [cut task failure](#)

rates by 9.5%. The Pro tier runs 100 tasks per day with 15 concurrent; Ultra runs 300 per day with 60 concurrent. The free tier on Gemini 3 Flash gives 15 per day.

The asynchronous model changes the unit of work. With Cursor or Codex you read every diff, type the next instruction, repeat. With Jules you fire off a task and walk away. When the resulting PR fails CI on GitHub Actions, Jules's [CI Fixer](#) reads the logs, pushes a fix to the same PR, and re-runs the checks on its own. Scheduled tasks handle weekly dependency bumps and nightly lint fixes on a cron. Sixty parallel tasks on Ultra runs like a small team of junior contributors filing PRs all day.

The interaction model is the bet. Cursor and Claude Code train you to think in tight prompt loops; Google trains you to think in pull requests. If the reviewer-only workflow scales, the bottleneck shifts from how fast you can type to how fast you can read and approve.

Read More: For Google's other I/O 2026 coding play, see [Antigravity 2.0 goes agent-only](#).

Sources:

- [Jules: Google's autonomous AI coding agent \(blog.google\)](#)
- [Jules Changelog](#)
- [Jules Getting Started Documentation](#)
- [Jules Product Page and Plans](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)