

# Kimi K2.7 Code: An Open 1-Trillion-Parameter Coding Model at a Fifth of Opus 4.8's Price

Kabui, Charles

2026-06-29

---

[Read at ToKnow.ai](#)

---

**Kimi K2.7 Code:**  
**An Open 1T Model**  
**at a Fifth the Cost**

Moonshot AI's model rivals Opus 4.8 on agent tool use

|                                                    |                                                  |                                                     |
|----------------------------------------------------|--------------------------------------------------|-----------------------------------------------------|
| <b>1T</b><br>Parameters, only 32B active per token | <b>5x</b><br>Cheaper to run than Claude Opus 4.8 | <b>256K</b><br>Token context, fits a whole codebase |
|----------------------------------------------------|--------------------------------------------------|-----------------------------------------------------|

Open weights under a Modified MIT license, free to self-host

June 29, 2026

ToKnow.ai

Moonshot AI released Kimi K2.7 Code on June 12, 2026, an open-weight coding model built on its earlier K2.6. It uses a mixture-of-experts design, so it holds 1 trillion parameters but activates only 32 billion per token to stay fast and cheap. It reads up to 256,000 tokens at once,

enough for a large codebase, and takes text, images, and video. It always works in a thinking mode that keeps its reasoning visible, yet cuts that thinking-token use by about 30% versus K2.6 to stop overthinking simple steps. Moonshot says it runs more than 12 hours nonstop, and on MCP Mark, which tests calling external tools, it scores 81.1, ahead of Claude Opus 4.8's 76.4.

It does not top every chart. It still trails GPT-5.5 and Opus 4.8 on most coding benchmarks, and GPT-5.5 scores 92.9 on that same test. Most of that capability comes much cheaper. Moonshot's API charges \$0.95 per million input tokens and \$4 per million output, against \$5 and \$25 for Opus 4.8, roughly 5 times cheaper to read and 6 times cheaper to write. The Modified MIT license also lets a team download the weights and self-host, or rent it through OpenRouter or Cloudflare Workers AI.

It lands the same month as Z.ai's [GLM-5.2](#), another open-weight coding model chasing the frontier. While OpenAI and Anthropic fight over the top score, Chinese labs keep shipping downloadable models that are close enough on agent work and far cheaper, moving the contest from raw capability toward price and ownership.

Read More: [Qwen3.6 and DeepSeek V4, China's open-weight models matching the frontier](#)

Sources:

- [Here comes Kimi K2.7 Code: Better Coding with more efficiency \(Moonshot AI forum\)](#)
- [Kimi-K2.7-Code model card and weights \(Hugging Face\)](#)
- [Kimi K2.7 Code pricing and providers \(OpenRouter\)](#)
- [kimi-k2.7-code on Cloudflare Workers AI](#)
- [Claude Opus 4.8 token pricing \(Anthropic docs\)](#)

---

**Disclaimer:** For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)