

A 37-Author Survey Says Latent Space, Not Language, Is Where AI Should Think

Kabui, Charles

2026-04-04

[Read at ToKnow.ai](#)

A 37-Author Survey Says Latent Space Is Where AI Should Think

Continuous representations as the native computational substrate

37 Authors from 16 institutions Multi-institutional survey	7 Ability domains mapped Reasoning to embodiment	511+ GitHub stars on reading list Curated paper collection
---	---	---

April 4, 2026

ToKnow.ai

A team of 37 researchers from 16 institutions published “[The Latent Space](#),” a survey arguing that continuous internal representations, not discrete text tokens, are becoming the primary medium where AI models actually compute. The paper organizes a fragmented field along two axes: mechanism (how latent spaces are built, covering architecture, representation, computation, and optimization) and ability (what they enable, spanning seven domains: reasoning,

planning, modeling, perception, memory, collaboration, and embodiment). It traces the field from early prototypes like [COCONUT](#), which first showed that continuous thought vectors can encode multiple reasoning paths simultaneously, through a current “outbreak” phase where latent methods have spread from text-only reasoning into visual cognition, robotic control, and multi-agent coordination. The [companion reading list](#) already has over 500 GitHub stars.

Standard chain-of-thought reasoning forces a model to spell out every intermediate step in text, adding latency and wasting compute on tokens that carry no logical weight. Latent-space methods skip that bottleneck: models reason in continuous vectors internally and output only the final answer. Work covered in the survey shows this can compress reasoning chains by orders of magnitude while maintaining accuracy. For anyone building reasoning systems, the survey is a structured map of which latent techniques work, for which modalities, and under what optimization strategies.

AI may be shifting from “thinking in words” to “thinking in vectors.” If that trajectory continues, the next generation of foundation models won’t treat latent space as a hidden implementation detail. It will be the primary workspace where reasoning, planning, and coordination happen, with language reserved for input and output.

Read More: [-Reasoner applies gradient descent directly in latent space during inference](#)

Sources:

- [The Latent Space: Foundation, Evolution, Mechanism, Ability, and Outlook \(arXiv\)](#)
- [Awesome-Latent-Space companion reading list \(GitHub\)](#)
- [COCONUT: Training Large Language Models to Reason in a Continuous Latent Space](#)
- [HuggingFace Daily Papers, April 3, 2026](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)