

Lock-Free Programming: Threads Share Data Without Waiting on a Lock

Kabui, Charles

2026-08-25

 Read at ToKnow.ai



**Lock-Free Programming:
Threads Share Data
Without Waiting on a Lock**

Atomic operations instead of mutexes, from Herlihy to LMAX

1,000x lower event latency vs queues	25M/s messages per second, under 50 ns	8x more throughput for the same setup
--	--	---

August 25, 2026 ToKnow.ai

Lock-free programming lets several threads share data without making them take turns on a mutex, the lock that normally pauses one thread while another holds it. Threads coordinate with atomic operations instead, most often compare-and-swap, a single indivisible step that updates a value only if it still matches an expected one. Because no thread ever blocks, a slow or paused thread can't stall everyone else. The idea traces to Herlihy's 1991 wait-free

synchronization paper, and the most used structure, a lock-free first-in-first-out queue, comes from [Michael and Scott's 1996 paper](#). The hard part isn't the atomic step; it's reclaiming memory safely and beating the ABA problem, where a value changes and then changes back.

Locks get expensive exactly where they're needed most. Under heavy contention, waiting on a lock means kernel arbitration and cache misses that can cost more than the work itself. At LMAX, a financial exchange, queues between processing stages added as much latency as disk I/O, hundreds of microseconds per hop. Its lock-free [Disruptor](#) cut mean latency for a three-stage pipeline by three orders of magnitude, handled roughly eight times the throughput, and passed over 25 million messages per second at under 50 nanoseconds. The same priorities drive [NautilusTrader](#), an open-source Rust trading engine with a deterministic, event-driven architecture built for low-jitter execution.

None of this makes locks obsolete. Lock-free code is famously hard to reason about, wait-free structures are rarer still, and most production code is lock-free only with careful memory reclamation on top. But when mutexes become the bottleneck, there is a proven ladder above them: atomic operations, then lock-free queues and rings, then wait-free designs.

Read More: [Bloom filters, another elegant trick that powers systems at scale](#)

Sources:

- [Herlihy \(1991\): Wait-Free Synchronization \(PDF\)](#)
- [Michael & Scott \(1996\): Simple, Fast, and Practical Non-Blocking and Blocking Concurrent Queue Algorithms \(PDF\)](#)
- [LMAX Disruptor: High Performance Alternative to Bounded Queues \(Technical Paper\)](#)
- [Preshing: An Introduction to Lock-Free Programming](#)
- [NautilusTrader: Rust-Native Trading Engine \(GitHub\)](#)

***Disclaimer:** For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)*