


Microsoft VibeVoice: A Model That Turns Text Into 90-Minute, Four-Speaker Audio

Kabui, Charles

2026-06-17

[Read at ToKnow.ai](#)



The image is a promotional graphic for Microsoft VibeVoice. It features a dark blue background with a grid pattern. The Microsoft logo is in the top left, and 'VibeVoice' is in a large, stylized font. Below the title, it says '90-minute, four-speaker speech from text, in a single pass'. There are three main features highlighted in colored boxes: '90 min' (orange) with sub-points 'Audio in a single pass' and 'No stitching of clips'; '4 voices' (yellow) with sub-points 'Distinct speakers, one take' and 'Natural turn-taking'; and '80x' (blue) with sub-points 'Smaller audio tokens' and 'vs the Encodec codec'. At the bottom, it says 'An open model for long-form podcasts and audiobooks, run locally', 'June 17, 2026', and 'ToKnow.ai'.

Microsoft VibeVoice

90-minute, four-speaker speech from text, in a single pass

- 90 min**
Audio in a single pass
No stitching of clips
- 4 voices**
Distinct speakers, one take
Natural turn-taking
- 80x**
Smaller audio tokens
vs the Encodec codec

An open model for long-form podcasts and audiobooks, run locally

June 17, 2026

ToKnow.ai

Microsoft Research built [VibeVoice](#), a text-to-speech model that turns a written script into up to 90 minutes of continuous audio with as many as four distinct speakers, in a single pass. Most open speech models manage one or two speakers for a couple of minutes; VibeVoice keeps each voice steady across a full podcast. The trick is a speech tokenizer running at just 7.5 frames per second, about 80 times more efficient than the common Encodec codec, which lets

a 64,000-token context hold an hour and a half of speech. A language model (Qwen2.5, in 1.5B and 7B sizes) tracks the dialogue while a small diffusion head fills in the sound one chunk at a time, an approach called [next-token diffusion](#). In tests, the 7B version beat ElevenLabs and Google’s Gemini text-to-speech on how realistic listeners found the audio.

This makes long-form audio cheap to produce. Hand VibeVoice a full podcast script and it returns an hour of natural back-and-forth between named speakers, instead of recording hosts or paying a commercial service per character. The model runs locally on a single GPU, with no subscription. For audiobooks, training material, and multi-host shows, a studio session becomes one generation step.

The enabling idea is compression. By squeezing each second of audio into just 7.5 tokens, VibeVoice fits 90 minutes into a context window most models spend on a few minutes. It points to a pattern in generative audio: better tokenizers, not just bigger models, unlock long-form output.

Read More: [Google DeepMind Lyria 3](#) takes generative audio in a different direction, producing full music tracks from a text prompt.

Sources:

- [VibeVoice Technical Report \(arXiv\)](#)
- [VibeVoice-1.5B model card \(Hugging Face\)](#)
- [VibeVoice on GitHub](#)
- [Multimodal Latent Language Modeling with Next-Token Diffusion \(arXiv\)](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)