MiniMax-M2.5: Frontier-Class Agentic Model at One-Tenth the Cost

Kabui, Charles

2026-02-20

Read at ToKnow.ai

Figure 1: MiniMax M2.5

MiniMax released M2.5, a 229B-parameter model trained with reinforcement learning across hundreds of thousands of real-world environments. It scores 80.2% on SWE-bench Verified (real-world bug fixing), 51.3% on Multi-SWE-bench (multilingual issue resolving), and 76.3%

on BrowseComp (complex web search tasks). On SWE-bench using the Droid scaffold, M2.5 scores 79.7%, beating Claude Opus 4.6's 78.9%. The model was trained using Forge, an agent-native RL framework that decouples the training engine from the agent scaffold, alongside a process reward mechanism for end-to-end monitoring. MiniMax reports that 80% of newly committed code inside their own company is now generated by M2.5, and 30% of all internal tasks across R&D, sales, HR, and finance are completed autonomously by the model.

The cost figures are the standout detail. M2.5-Lightning runs at 100 tokens per second and costs $0.3 per million input tokens, $2.4 per million output tokens. Running it continuously for one hour costs $1. At the slower 50 TPS tier, that drops to $0.30 per hour. MiniMax states this is one-tenth to one-twentieth the cost of Opus, Gemini 3 Pro, and GPT-5. The model completes SWE-bench tasks in 22.8 minutes on average, matching Opus 4.6's 22.9 minutes, at a fraction of the price.

When a model matches frontier performance at 5-10% of the cost, the competitive dynamics shift from capability to economics. Startups and independent developers can now run agent workloads that previously required enterprise-tier API budgets. The open weights (available on Hugging Face under a modified MIT license) mean self-hosting is viable too.

Sources:

- MiniMax-M2.5 on Hugging Face
- MiniMax API Platform
- MiniMax Agent

Disclaimer: *For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation.* ***Read more:*** */terms-of-service*