OpenAI's Agents API: One Call, and OpenAI Runs the Agent

Kabui, Charles

2026-09-17

Read at ToKnow.ai



Think of an AI model as a brain. Real work also needs the parts around it: memory that survives long jobs, hands to run code and use tools, and a workspace. OpenAI built all of that for its coding agent, Codex, and on September 10, 2026 started renting it out through the Agents API. A developer names a model, some tools, and where to run it, then makes one call. OpenAI keeps the session alive for hours, summarizes old context before it overflows,

spreads work across up to three helper agents, and loads tool definitions only when needed. Agents run in an OpenAI sandbox, on your own machines, or through nine partners including Cloudflare, Modal, and Vercel. You pay standard model, tool, and container rates, with no platform fee.

It is plumbing for developers, not a chatbot you open. Anyone building an agent that runs unattended assembles a job queue, a state database, a sandbox fleet, a compaction routine, and a retry policy, each with an owner. Renting that removes months of setup. The cost is dependence: model, memory, tools, and workspace all come from one vendor, and an analyst told InfoWorld it does not support zero data retention even on your own sandbox, which limits use in regulated industries.

OpenAI is not alone here. Anthropic's Claude Managed Agents has been in public beta since April, and AWS's Bedrock AgentCore went generally available in June, letting teams switch models mid-session. OpenAI's version bets teams trade flexibility for less work.

Read More: OpenAI Codex CLI: A Rust-Built Open-Source Coding Agent

Sources:

- Introducing the Agents API
- Agents API overview
- Codex harness on GitHub
- OpenAI API pricing
- InfoWorld: OpenAI launches managed Agents API

Disclaimer: *For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation.* ***Read more:*** */terms-of-service*