

Protea: A South African-Founded Startup Ships a 10 Million-Token Context Window

Kabui, Charles

2026-08-26

 Read at [ToKnow.ai](#)



The graphic features a dark blue background with a subtle grid and a network diagram in the top right. The main title is in large white and green font. Below it, a subtitle states 'Up to 10 million tokens in a single pass, no waitlist'. Three key metrics are highlighted in separate boxes: '10M tokens in a single context window', '7.5M words, about 15,000 pages', and '95-99% accuracy kept after compression'. The date 'August 26, 2026' and the 'ToKnow.ai' logo are at the bottom.

Protea: A South African-Founded Startup Ships a 10 Million-Token Context Window

Up to 10 million tokens in a single pass, no waitlist

10M tokens in a single context window	7.5M words, about 15,000 pages	95-99% accuracy kept after compression
---	--	--

August 26, 2026

ToKnow.ai

Refiant, a startup founded in South Africa, launched Protea, a family of long-context models that hold up to 10 million tokens in a single pass. That is roughly 7.5 million words, about 15,000 pages, or five years of email and Slack messages. Protea is available now through an API or chat interface in 1 million, 5 million, and 10 million token windows, with no waitlist. Refiant was founded in 2025 by a quantum mathematician and two co-founders, and it builds

on a compression technique that already shrank a 120-billion-parameter model to 12 GB of RAM while keeping 95-99% of its accuracy.

Whole codebases, regulatory archives, or years of claims data can now be fed to a model in one session instead of being chopped into chunks for a retrieval system. The South African founding is part of the story too: a model family at this scale shipping from a team based in Durban and Cape Town, with research partners at Imperial College London and UCL. Refiant's own tests claim the compressed models are up to 100 times more energy efficient than data center hardware, though these are self-reported numbers and long-context models still lose accuracy on material buried in the middle of a window.

Context size is turning into a shipping race, not just a research one. Magic announced a 100-million-token model in 2024 and nothing public shipped, a pattern we covered with [SubQ](#). Refiant put 10 million tokens in a free API within a year of founding. Efficiency, not just scale, is how a small team competes with labs spending billions.

Read More: [AI memory: why bigger context windows don't stop agents from forgetting](#)

Sources:

- [Refiant Launches 10 Million Token Long-Context Window \(Refiant\)](#)
- [Refiant Benchmark Results \(Refiant\)](#)
- [AI Startup Refiant Raises \\$5M to Slash Energy Footprint of Artificial Intelligence \(Refiant\)](#)
- [South African-Founded Refiant Launches 10 Million Token Context Window \(TechFinancials\)](#)
- [Refiant Unveils 10 Million-Token AI Model \(BusinessDay\)](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)