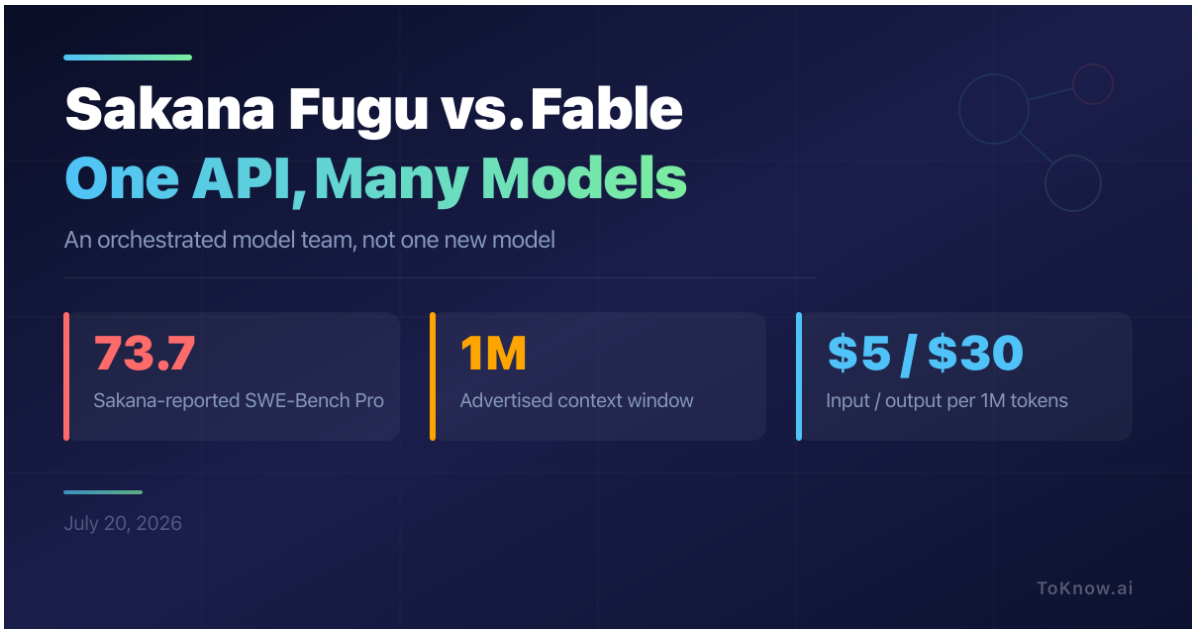
Sakana Fugu vs. Fable: A Model Team Behind One API

Kabui, Charles

2026-07-20

Read at ToKnow.ai



Sakana AI's Fugu is not one new foundation model. It is a trained coordinator that builds a workflow for each request, choosing workers, assigning roles, and combining their answers behind one API. The technical report names Gemini 3.1 Pro, Claude Opus 4.8, and GPT-5.5 among its worker pool. Fable 5 and Mythos Preview are comparison models, not hidden Fugu workers. Sakana says neither is publicly accessible. Its own tests put Fugu Ultra at 73.7 on

SWE-Bench Pro and 82.1 on Terminal-Bench 2.1. Sakana ran Fugu but took the baseline scores from model providers, so the table is not one independently reproduced test.

Existing OpenAI-compatible clients can call Fugu without managing several APIs. Fugu Ultra advertises a one-million-token context window and costs $5 for input and $30 for output per million tokens. Rates rise to $10 and $45 when the input context exceeds 272,000 tokens. The closed routing is the catch. Sakana does not reveal which workers handled a request, may use customer data for training unless the user opts out, and does not serve the EU or EEA. Its public repository contains setup files, demos, and the report, but no coordinator weights or software licence. That makes Fugu a closed service built from multiple models, not an open model that replaced Fable.

Read More: when one self-managing agent is simpler than a multi-agent stack.

Sources:

- Sakana AI: Fugu announcement, pricing, and FAQ
- Sakana Fugu Technical Report
- SakanaAI/fugu repository
- OpenRouter: Fugu Ultra pricing and context

Disclaimer: *For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation.* ***Read more:*** */terms-of-service*