

Seoul World Model: The First AI That Simulates a Real City

Kabui, Charles

2026-03-18

[Read at ToKnow.ai](#)



The graphic features a dark blue background with a subtle grid pattern. In the top right corner, there is a faint, stylized line graph with three data points. The main title 'Seoul World Model' is in large white font, followed by 'The First AI That Simulates a Real City' where 'City' is highlighted in green. Below the title, a subtitle reads 'City-scale world simulation grounded in actual street-view data'. Three vertical bars in orange, yellow, and blue separate three columns of data. The first column shows '1.2M' in orange, with 'Real street-view panoramas' and 'Anchoring generation to Seoul' below it. The second column shows '3 Cities' in yellow, with 'Seoul, Busan, Ann Arbor' and 'Validated across real metropolises' below it. The third column shows '10K' in blue, with 'Synthetic training videos' and 'From Unreal Engine simulator' below it. At the bottom left, the date 'March 18, 2026' is displayed, and at the bottom right, the 'ToKnow.ai' logo is visible.

Seoul World Model

The First AI That Simulates a Real City

City-scale world simulation grounded in actual street-view data

1.2M Real street-view panoramas Anchoring generation to Seoul	3 Cities Seoul, Busan, Ann Arbor Validated across real metropolises	10K Synthetic training videos From Unreal Engine simulator
--	--	---

March 18, 2026

ToKnow.ai

NAVER AI Lab and KAIST released [Seoul World Model](#) (SWM), the first city-scale world model that generates video of a real city rather than an imagined one. Given a location, camera path, and optional text prompt, SWM retrieves nearby images from [1.2 million street-view panoramas](#) of Seoul and uses them to anchor each video segment to actual buildings, roads, and landmarks. Cross-temporal pairing forces the model to learn permanent structure

rather than copy transient objects like parked cars. A Virtual Lookahead Sink re-anchors each segment to a reference image further ahead on the route, preventing quality drift over long distances. Tested on [three cities](#), SWM outperforms existing video world models in spatial fidelity and consistency over trajectories of hundreds of meters.

Previous world models produce convincing but fictional environments. SWM generates video of streets that actually exist. Autonomous driving teams can test routes through virtual Seoul before real roads. Urban planners can visualize changes to real neighborhoods. Training combines real panoramas with [10,000 synthetic videos from Unreal Engine](#) covering 431,500 square meters. You can even reshape real scenes with text prompts: summon a tidal wave onto Gangnam or drop Godzilla between actual skyscrapers.

Most generative video research optimizes for visual realism in fictional worlds. SWM reverses the priority: anchor to reality first, then generate. Retrieval over real geospatial data is emerging as the bridge between generative AI and digital twins of actual cities.

Read more: [World Labs' \\$1B bet on spatial AI world models](#)

Sources:

- [Seoul World Model Paper \(arXiv\)](#)
- [Seoul World Model Project Page](#)
- [GitHub Repository \(310 stars\)](#)
- [HuggingFace Daily Papers \(118 upvotes\)](#)

***Disclaimer:** For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)*