


Supertonic 3: 99M-Parameter On-Device TTS Across 31 Languages, No GPU Required

Kabui, Charles

2026-05-28

[Read at ToKnow.ai](#)



Supertonic 3: On-Device TTS Across 31 Languages, No GPU Required

99M-parameter open-weight model via ONNX Runtime

99M Parameters vs 0.7B-2B for competing TTS	31 Languages from a single checkpoint	11 Platform SDKs from Python to Flutter
---	---	---

May 28, 2026

ToKnow.ai

Supertone released [Supertonic 3](#), a 99-million-parameter text-to-speech model that runs entirely on-device through ONNX Runtime (a cross-platform inference engine). It supports 31 languages from a single checkpoint, outputs audio at 44.1kHz, and needs no GPU. The model runs on desktop, mobile, Raspberry Pi, and even e-readers in airplane mode. Ten inline expression tags (<laugh>, <breath>, <sigh>) add human nuance without reference

audio. A language-agnostic mode auto-detects input language when unspecified. SDKs cover [11 platforms](#) including Python, Node.js, WebGPU browsers, Java, C++, Swift, Rust, and Flutter. The [Python SDK](#) includes a local HTTP server with an OpenAI-compatible `/v1/audio/speech` endpoint, so existing apps can switch by changing a URL. Weights are open under the [OpenRAIL-M license](#).

Cloud TTS services charge per character and require an internet connection. Supertonic 3 runs offline, costs nothing after download, and fits in 99M parameters, a fraction of the 0.7B to 2B range typical of competing open TTS systems. A developer building a voice assistant can run `pip install supertonic` and generate speech locally in under a second. An e-reader can narrate books with zero network dependency. Any tool already using the OpenAI TTS API can switch to Supertonic's local server with one URL change.

On-device audio AI has closed the gap with cloud services faster than most expected. When a 99M-parameter model produces usable quality across 31 languages on a \$35 Raspberry Pi, the argument for paying per-character API fees weakens for most use cases outside premium voice cloning. For more on the broader on-device AI trend, see [Unslloth Studio + Qwen3.6: Running Frontier AI on 18GB RAM](#).

Sources:

- [Supertonic 3 GitHub Repository](#)
- [Supertonic 3 Model on HuggingFace](#)
- [Supertonic 3 Audio Samples and Demo Page](#)
- [Supertonic Python SDK Documentation](#)
- [Supertonic 3 Interactive Demo on HuggingFace Spaces](#)

Disclaimer: For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)