

# Vinci2 Tests When a Camera Assistant Should Speak Up

Kabui, Charles

2026-07-26

---

[Read at ToKnow.ai](#)

---



Vinci2 is a research system for camera assistants that must decide whether to speak without being asked. Its [EgoServe benchmark](#) marks 3,437 useful interventions across 128.5 hours of first-person video, split among 10 service types and four time ranges. These run from an immediate safety warning to advice based on habits seen days earlier. The training-free EgoMemo agent turns incoming video into layered summaries, a graph of people and objects,

and a searchable visual archive. At each step, it retrieves relevant history before choosing whether to respond. That memory helped, but the [reported results](#) are still modest: EgoMemo reached 8.0 overall F1, a measure balancing correctly timed help against missed or misplaced interventions, compared with 4.7 for GPT-5-mini and 3.5 for Qwen3-VL-Plus.

A wearable assistant could warn someone about unsafe knife handling, recover a missed recipe step, or recall an unfinished task without waiting for a prompt. EgoServe gives researchers a way to test those moments rather than grading answers to prewritten questions. The [MIT-licensed code](#) and CC BY 4.0 annotations are public, but the source videos require separate access. The release also does not establish private, on-device operation: its setup asks for several cloud API keys. More importantly, the [paper](#) reports no user study of interruption tolerance, battery cost, or bystander privacy. Before this belongs in everyday glasses, staying silent needs to become a measured safety behavior, not just another model output.

Sources:

- [Vinci2 paper](#)
- [Vinci2 project page and results](#)
- [EgoMemo code and evaluation suite](#)
- [EgoServe dataset card](#)
- [EgoLife source dataset](#)

---

***Disclaimer:** For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)*