


Alibaba's Wan3.0 Turns Documents Into 30-Second Videos in One Pass

Kabui, Charles

2026-08-25

[Read at ToKnow.ai](#)



The graphic features a dark blue background with a subtle grid and a faint line-art illustration of a person in the upper right. The main title is in large white and green font. Below it, a subtitle in light blue text states 'Up to 30 seconds of footage from text, images, audio, video, or documents'. Three key metrics are highlighted in separate boxes with vertical bars: '30s video in a single pass' (red bar), '\$6.00 full 30s 1080P clip via API' (yellow bar), and '\$0.05/s 480P API price per second' (cyan bar). At the bottom left, a date 'August 25, 2026' is shown next to a small green progress bar. The 'ToKnow.ai' logo is in the bottom right corner.

Alibaba's Wan3.0 Turns Documents Into 30-Second Videos in One Pass

Up to 30 seconds of footage from text, images, audio, video, or documents

30s video in a single pass	\$6.00 full 30s 1080P clip via API	\$0.05/s 480P API price per second
--------------------------------------	--	--

August 25, 2026

ToKnow.ai

Alibaba officially [rolled out Wan3.0 on August 24](#), after a public beta that began August 6. It is a video generation model that produces up to 30 seconds of footage in a single pass, and the input side is the story: alongside text, images, audio, and video, it accepts [office documents](#), doc, xls, ppt, pdf, and markdown files, plus web page links, one file up to 100 MB or 50 pages per request. It runs on Alibaba Cloud Model Studio, with API pricing of \$0.05 per second

at 480P, \$0.10 at 720P, and \$0.20 at 1080P, so a full 30-second 1080P clip comes to \$6.00. Alibaba Cloud is [discounting API prices 30% until September 23](#).

A sales deck becomes a product launch ad, a training document becomes video courseware, a spreadsheet becomes animated charts, and a quarterly report becomes a narrated briefing, with no scriptwriter or editor in between. That collapses the usual multi-tool pipeline into one API call, and at \$1.50 per 30-second 480P draft, iteration is cheap. Wan3.0 also lets you edit instead of regenerate: the built-in editor can change visuals, plot, or dialogue on an existing clip. On-screen text rendering and audio texture are still weak spots.

Video generation is moving from short clips to anything-to-video, and when the source can be the documents people already make for work, generated video becomes the default way to present ordinary office content.

Read More: [ByteDance Seedance: one model that takes text, images, audio, and video together to make clips](#)

Sources:

- [Wan3.0: 30-Second AI Video Generation from Any Input \(Alibaba Cloud\)](#)
- [Alibaba launches Wan3.0 AI video model after \\$10 billion share sale \(Reuters\)](#)
- [Alibaba launches Wan3.0 video model with 30-second generation and document input \(TechNode\)](#)
- [Wan 3.0 AI Video Model: 30-Second One-Take, Document-to-Video \(PicLumen\)](#)

***Disclaimer:** For information only. Accuracy or completeness not guaranteed. Illegal use prohibited. Not professional advice or solicitation. **Read more:** [/terms-of-service](#)*